

Experiment 5 Postlab:

1. Explain the Time Complexity of the A* Algorithm.

Ans: The time complexity of A* depends on the quality of the heuristic function. In a worst-case, the algorithm can be $O(b^d)$, where b is the branching factor – the average number of edges from each node, and d is the number of nodes on the resulting path.

2. What are the limitations of A* Algorithm?

Ans: The A* algorithm, while renowned for its efficiency, has certain limitations to consider:

1. **Computational Cost:** A* can be computationally expensive, especially in scenarios with:

Extensive search spaces: When dealing with a vast number of possible paths, the exploration process can become resource-intensive.

High branching factors: If each node in the search space has many potential neighbors, the algorithm needs to evaluate numerous options, increasing computation time.

2. *Reliance on Heuristics:*

A* heavily relies on the quality of the heuristic function used to estimate the distance to the goal.

Poor heuristic: A poorly designed heuristic can lead the algorithm down suboptimal paths, potentially missing the truly optimal solution.

Domain-specific: Designing an effective heuristic often requires significant domain knowledge, making it less versatile for problems with diverse characteristics.

3. *Limited Applicability:*

A* may struggle with specific types of search spaces:

Dynamic environments: A* assumes a static environment where the cost of moving between nodes remains constant. This can be problematic in real-world scenarios with dynamic elements.

Uninformed search spaces: If there is no information available to guide the search (i.e., no heuristic function), A* loses its efficiency advantage and might not be the most suitable choice.

3. Discuss A*, BFS, DFS, and Dijkstra's algorithm in detail with examples.

Ans:

Breadth-First Search (BFS):

BFS is an algorithm for traversing or searching tree or graph data structures.


It starts at the root (or an arbitrary node) and explores all of the neighbor nodes at the present depth before moving on to the nodes at the next depth level.

BFS uses a queue data structure to keep track of the next nodes to visit.

BFS guarantees the shortest path in terms of the number of edges from the start node to every other node.

Example:

Let's say we have a graph represented as follows:



If we perform BFS starting from node A, the traversal order would be: A -> B -> C -> D -> E -> F -> G -> H.

Depth-First Search (DFS):

DFS is another algorithm for traversing or searching tree or graph data structures.

It starts at the root (or an arbitrary node), and explores as far as possible along each branch before backtracking.

DFS uses a stack data structure (or recursion) to keep track of the next nodes to visit.

DFS does not guarantee the shortest path; it might traverse deep into a branch before considering other branches.

Example:

Using the same graph as above, if we perform DFS starting from node A, the traversal order might be: A -> B -> E -> F -> C -> D -> G -> H.

Dijkstra's Algorithm:

Dijkstra's algorithm is used to find the shortest path from a single source vertex to all other vertices in a weighted graph with non-negative edge weights.


It maintains a set of vertices whose shortest distance from the source vertex is already known.

It iteratively selects the vertex with the smallest tentative distance from the source vertex and updates the distances to its adjacent vertices if a shorter path is found.

Dijkstra's algorithm can be implemented using a priority queue.

Example:

Consider the following graph with weighted edges:



If we apply Dijkstra's algorithm with source node A, it would compute the shortest paths to all other nodes in the graph.

A Algorithm:*

A* (pronounced "A star") is a widely used algorithm for finding the shortest path in a graph between a given start node and a target node.

It uses a heuristic to estimate the cost from the current node to the goal, combined with the actual cost of reaching that node from the start.

A* keeps track of the total estimated cost of reaching a node, called the "f-score," which is the sum of the cost from the start to the current node (g-score) and the heuristic estimate from the current node to the goal (h-score).

It prioritizes exploring nodes with lower f-scores, thus guiding the search towards the goal more efficiently.

Example:

Suppose we have a grid-based map where each cell represents a node in the graph, and we want to find the shortest path from the start cell to the target cell. A* algorithm can be applied, where the heuristic function might be the Euclidean distance between the current cell and the target cell.